

Figure S3-2. Study 3: LLM-as-a-judge vs the human ceiling on 3 datasets (n=500 each, 6-model panel)

ChaosNLI-SNLI
(NLI, ceiling=low)

ChaosNLI-MNLI
(NLI, harder)

HateXplain
(hate speech, sensitive)

bar colour = reliability verdict; whiskers = 95% bootstrap CI; WR = alt-test winrate

