

Figure S3-2. Study 3: LLM-as-a-judge vs the human ceiling on 3 datasets (n=500 each, 6-model panel)
bar colour = reliability verdict; whiskers = 95% bootstrap CI; WR = alt-test winrate

