

When Can an Open LLM-as-a-Judge Replace Human Evaluation? A Cross-Domain Review and Empirical Test on Summarization and Natural Language Inference

ANONYMOUS AUTHOR(S)

Large language models are increasingly used as “judges” that stand in for human annotators, but the field lacks an evidence-based account of *when* an LLM judge agrees with humans well enough to substitute for them, and *what* governs that agreement. We address this for *open, locally-servable* models, the realistic and reproducible choice for most researchers. We contribute three studies anchored to a single principle: an LLM judge should be measured against the *human-human agreement ceiling* for a task, never a manufactured “gold.” **Study 1** is a PRISMA systematic review of 215 studies across 29 domains (all 215+55 references independently web-verified) that maps where LLM-as-a-judge agreement reaches the human ceiling. **Study 2** replicates the standard summary-level G-Eval protocol on SummEval with the most powerful open models and finds, with bootstrap confidence intervals, that open judges meet or exceed the established proprietary-judge correlations on the *verifiable* aspects (consistency, coherence) while the most compressed aspect (fluency) is hard to rank for every judge. **Study 3** is a new meta-evaluation of a six-model open panel (9B-754B) on natural-language-inference datasets with rich human label distributions (~100 annotators/item; 6,000 judgments), benchmarked against the annotator-bootstrap human ceiling with chance-corrected agreement, jury aggregation, a calibrated contamination screen, and an item-level mixed-effects analysis. The organizing finding is that judge reliability is governed by *item-level human disagreement*: panel accuracy on low-disagreement items (0.79) far exceeds that on high-disagreement items (0.61), item-level variance dominates dataset- and model-level variance, and model scale contributes only modestly. We turn this into a decision framework and an eight-point protocol for designing LLM-as-a-judge experiments. A companion paper applies the same evaluation, with a systematic review, to subjective social-computing annotation in Human-Centered Computing.

Additional Key Words and Phrases: LLM-as-a-judge, automatic evaluation, human-agreement ceiling, annotator disagreement, summarization, natural language inference, reproducibility

1 INTRODUCTION

Large language models are now routinely deployed as *judges*: instead of recruiting human annotators to label data, rate outputs, or score quality, researchers prompt an LLM to do it for a fraction of the cost [15, 23, 32]. The practice is spreading faster than the evidence that licenses it. Two questions remain under-answered. **When** does an LLM judge agree with humans closely enough to substitute for them, and **what** governs that agreement, model scale, task, or something about the human labels themselves?

We answer these questions for *open, locally-servable* models. This is a deliberate scope, not a limitation: proprietary judges are expensive, non-reproducible, and change without notice, so for most researchers the realistic question is whether an open model they can run themselves is a trustworthy judge. Our entire empirical program uses open models served through a standard vLLM endpoint at temperature 0 with released harnesses.

One methodological commitment runs through the paper: an LLM judge must be measured against the *human-human agreement ceiling* for the task, not against a single “gold” label. For subjective constructs there is no single correct answer [2, 25]; the relevant question is whether the judge agrees with humans as well as humans agree with each other. We compute that ceiling directly from the per-item annotator distributions that the datasets already provide, so our human reference is real, independent, and inherited from the data, not added by us.

Contributions. (1) **Study 1**: a PRISMA systematic review of 215 studies across 29 domains (every reference independently web-verified) that maps, against the human ceiling, where LLM-as-a-judge is reliable and where it is not. (2)

Study 2: a controlled replication showing that the most powerful open models now meet or exceed the established proprietary-judge correlations on the verifiable aspects of a flagship benchmark, with bootstrap confidence intervals. (3) **Study 3:** a new meta-evaluation of a six-model open panel (9B-754B) on NLI datasets with rich human label distributions (6,000 judgments), establishing that *item-level human disagreement*, not model scale or data contamination, is the binding constraint on judge reliability. (4) A decision framework and an eight-point protocol for designing LLM-as-a-judge experiments. All code, corpora, search logs, and harnesses are released.

2 BACKGROUND AND SHARED METHODOLOGY

LLM-as-a-judge descends from reference-free metrics (BERTScore, COMET) but uses a general instruction-following model with a natural-language rubric, supporting pointwise, pairwise, and listwise judgment [16, 22]. Four architectural families recur: prompted general LLMs [23, 32]; fine-tuned open evaluators [19, 28, 33]; juries/multi-agent panels [8, 27]; and calibrated rubric methods [17]. Prior surveys taxonomize methodology [16, 21]; we instead synthesize *where substitution is empirically validated* and quantify it against the human-human ceiling.

Related work. Four literatures meet in this paper. (1) *LLM-as-a-judge evaluation*: studies report strong agreement on aggregate, verifiable tasks (MT-Bench [32], GEMBA [20], SAFE factuality [30]) but document systematic biases (position, verbosity, self-preference, and an expertise gap by which judges converge on crowd rather than expert consensus [3]), which motivates benchmarking against the human ceiling rather than against a single “gold.” (2) *LLMs in HCI and qualitative research*: a fast-growing body uses LLMs to assist coding and thematic analysis [11, 14, 31] while qualitative scholars warn that interpretation, reflexivity, and positionality resist automation [12, 26]. (3) *Participant and persona simulation*: “silicon samples” [1] and synthetic users are increasingly scrutinized for flattening identity and distorting variance [5, 9, 29]. (4) *Learning with disagreement and perspectivism*: work on annotator disagreement and human label variation [6] argues that for subjective constructs the “ceiling” is itself low and contested, which our Study 3 datasets (ChaosNLI) operationalize directly. We connect these threads with two methodological commitments often missing from judge evaluations: anchoring every verdict to the *human-human* ceiling, and probing *data contamination* so that agreement on older public datasets is not mistaken for genuine competence.

Shared protocol. The systematic review follows PRISMA 2020 [24], identifying studies through a triangulated search (the OpenAlex API as the primary bibliographic database, Semantic Scholar as a cross-check, and a structured web “deep-research” arm), and code each study on a five-level **reliability verdict** benchmarked against the *human-human agreement ceiling* for the same task: **1 Validated** ($\approx$ / $>$ ceiling), **2 Promising** (conditional), **3 Mixed** (task-dependent), **4 Caution** (below ceiling), **5 Unreliable** (near chance / bias-dominated). We report a reliability score = 6–verdict and synthesize within metric families rather than pooling incommensurable effect sizes. Full protocol, search logs, and the coded corpus are released.

3 STUDY 1: THE CROSS-DOMAIN RELIABILITY LANDSCAPE

3.1 Method (summary)

Study 1 systematically reviewed 215 studies across 29 domains comparing LLM-as-a-judge to human annotation. Identification triangulated OpenAlex (1500 records, 957 unique), Semantic Scholar, and a 59-query web search; screening retained on-topic empirical comparisons; 29 studies reported an extractable numeric agreement value (Figure 1; the full coded corpus is Table 6). **Coding reliability:** because the verdict rubric embeds judgement, we validated it with

Figure 1. PRISMA 2020 flow - two-arm identification (OpenAlex database + web/citation)




Fig. 1. Study 1 PRISMA 2020 flow (two-arm identification: OpenAlex database + web/citation).

an independent re-coding of a 25% random sample ($n=54$): the reliability verdict reproduced at Krippendorff $\alpha = 0.87$ (ordinal) with 100% of cases within one level (Appendix D). The re-coders are LLM-assisted, like the primary pipeline, so this measures rubric *reproducibility* rather than human inter-rater agreement, which we flag as a limitation. **Citation integrity:** every one of the 215+55 references was independently web-verified for existence against arXiv/venue/DOI records; 270/272 were confirmed (260 at high confidence) and the 2 that could not be verified were removed, so the corpus contains no unverifiable citations.

3.2 Findings: a green/amber/red landscape

The modal verdict is *mixed/task-dependent*; only 31% of studies sit in the reliable band and 19% in the problematic band (mean reliability 3.15/5). Reliability is governed by three properties: **objectivity**, **tractable verification**, and a **low human-agreement floor**. **Green (≥ 3.5 ; 8 domains):** machine translation (system level), text-to-image alignment, verifiable instruction-following, QA/factuality, information extraction, general-NLG pairwise preference, dialogue, and legal scoring; here LLM judges frequently match or exceed the human ceiling (e.g., MT-Bench 85% vs. 81% human-human [32]; GEMBA-MQM 96.5% [20]; SAFE wins 76% of disagreements [30]). **Amber (19 domains):** summarization, RAG, education, social-science annotation, medicine (structure-dependent: $\kappa \approx 0.82$ -0.93 structured vs. 0.47-0.71 complex reasoning), code, and IR/relevance (strong but contested [10]). **Red (2 domains):** low-resource multilingual evaluation




Fig. 2. Study 1. (a) Mean reliability by domain (green/amber/red). (b) Reliability landscape: domain×judging task; verification/classification columns are greener than open-ended scoring.

and creative writing (near-zero expert correlation [7]). Figures 2a-2b give the domain ordering and the domain×task landscape; the decision map (Figure 3b) positions domains by reliability against stakes.

3.3 Failure modes and the decision map

Recurring biases such as verbosity/length, position, self-preference, style-over-substance, preference leakage, non-determinism, and the *expertise gap* (judges converge to crowd, not expert, consensus [3]) all modulate reliability (Figure 3a). Agreement should be benchmarked against the human ceiling, which on subjective tasks is itself low (Krippendorff $\alpha < 0.35$ [18]). The Study-1a decision map (Figure 3b) and Table 1 translate this into deployment guidance and motivate the moving-target test (Study 2) that follows.

4 STUDY 2: IS THE RELIABILITY LANDSCAPE A MOVING TARGET?

A systematic review necessarily captures results *as reported*, and most “promising”-band evidence in Study 1 was produced with proprietary judges that are now one or two model generations old (e.g. the influential G-Eval used GPT-4). This raises a question that the review alone cannot answer: when an LLM judge falls short of the human ceiling in a promising domain, is that a *temporary* limitation that more capable models will erase (“matter of time”), or a *structural* property of LLM-as-a-judge that better models do not overcome? Study 2 answers this directly for a flagship case by re-running the evaluation with the most powerful open models available to us today and comparing against the canonical reported numbers.

4.1 Method

We select **SummEval** [13] (summarization quality assessment, a central amber-band task in Study 1) because it pairs (i) a widely cited LLM-judge result to beat (G-Eval/GPT-4; 23) with (ii) public expert human ratings on four aspects (coherence,




Fig. 3. Study 1. (a) Catalogue of reported failure modes/biases. (b) Deployment decision map (reliability $\times$ stakes; bubble area $\propto$ studies).

consistency, fluency, relevance) for $100 \times 16 = 1600$ machine summaries. We replicate the G-Eval summary-level protocol *verbatim* (same aspect definitions, 1-5 integer ratings, temperature 0) but swap the judge for the 4 most powerful open models in the available open-model catalogue (deepseek-v4-pro, glm-5.1, gptoss-120b, qwen3.6-35b-a3b). The comparison metric is the canonical *summary-level Spearman*: for each source document, correlate the judge’s 16 scores with the human means, averaged across documents. The bar to beat is G-Eval/GPT-4’s reported per-aspect correlations (coherence .582, consistency .507, fluency .455, relevance .547; average $\approx .52$). In total Study 2 contributes 25,600 fresh judgments.

4.2 Results

We follow the standard summary-level G-Eval protocol [23] and attach bootstrap 95% CIs (over the 100 source documents) to every correlation (Figure 4, Table 2). On the verifiable aspects the open judges *meet or exceed* the established G-Eval/GPT-4 benchmark: 5 of the 16 model $\times$ aspect cells have a 95% CI lying *entirely above* the reported value, all on **consistency** and **coherence** (e.g. deepseek-v4-pro consistency $p=0.69$, CI [0.63, 0.74] vs. .507), and **relevance** is on par (.48-.54 vs. .55). Our protocol differs from the original in two deliberate, deployment-relevant ways: the judges are open and locally-servable, and we read a single integer score rather than a proprietary model’s probability-weighted decoding. The comparison is to the figures reported by the original work, the standard reference point in this literature. But on **fluency** every open model falls well *short* (.22-.39 vs. a reported .455). We diagnose why: fluency is the most *compressed* aspect of the human reference, with the lowest within-document rating variance (SD = 0.63 vs 0.81 for consistency) and 80% of summaries rated near the top of the scale, so it offers the least to rank, and it is the aspect on which the *reported proprietary judge also scores lowest*. This is a property of the *reference*, not the judges: even a noiseless judge returning the rounded human means tops out at $p \approx 0.87$ on fluency. The clean result is therefore that capability gains concentrate on the verifiable, rankable aspects (consistency, coherence), while the most compressed aspect is hard to rank for every judge, open or proprietary.

Table 1. Domain-level decision table. Mean reliability score (1 = unreliable . . . 5 = validated), modal verdict, and traffic-light signal (GREEN ≥ 3.5 ; AMBER 2.5-3.5; RED < 2.5). Sorted by mean reliability. Small- n domains should be read together with the family-level aggregates.

Domain / task	n	Mean rel.	Modal verdict	Signal
Machine Translation	3	4.33	Promising	GREEN
Text-to-Image	2	4.00	Promising	GREEN
Dialogue	1	4.00	Promising	GREEN
Instruction Following	3	4.00	Validated	GREEN
QA & Factuality	5	3.80	Promising	GREEN
Information Extraction (NER/NLI)	9	3.78	Promising	GREEN
General NLG	21	3.71	Promising	GREEN
Legal	2	3.50	Promising	GREEN
Education & Grading	17	3.35	Mixed	AMBER
RAG	3	3.33	Mixed	AMBER
Summarization	8	3.25	Promising	AMBER
Fact-Checking	4	3.25	Promising	AMBER
Sentiment & Emotion	4	3.25	Mixed	AMBER
IR & Relevance	9	3.22	Promising	AMBER
Social-Science Annotation	15	3.20	Mixed	AMBER
Peer Review & Science	6	3.17	Mixed	AMBER
Multimodal (Vision-Language)	5	3.00	Mixed	AMBER
Qualitative Coding	6	3.00	Mixed	AMBER
Argument Quality	1	3.00	Mixed	AMBER
Finance	3	3.00	Mixed	AMBER
Medicine & Clinical	16	2.88	Mixed	AMBER
Methodology & Bias	38	2.87	Mixed	AMBER
Reward Modeling	5	2.80	Mixed	AMBER
Code & SW Eng	7	2.71	Mixed	AMBER
Reasoning & Math	3	2.67	Mixed	AMBER
Content Moderation & Toxicity	5	2.60	Mixed	AMBER
Mental Health	4	2.50	Mixed	AMBER
Creative Writing	6	2.33	Mixed	RED
Multilingual	4	2.25	Caution	RED

Table 2. Study 2. Summary-level Spearman ρ of the most powerful open judges on SummEval vs. the reported G-Eval-4/GPT-4 correlations (25,600 judgments). Green cells *surpass* the reported per-aspect number.

Judge	Coher.	Consist.	Fluency	Relev.	Avg ρ
<i>Reported (G-Eval-4)</i>	0.582	0.507	0.455	0.547	0.523
deepseek-v4-pro	0.548	0.688	0.392	0.540	0.542
qwen3.6-35b-a3b	0.606	0.702	0.298	0.540	0.536
glm-5.1	0.650	0.693	0.217	0.543	0.526
gptoss-120b	0.586	0.685	0.260	0.482	0.503

Interpretation: time vs. structure. Study 2 shows the reliability landscape is *partly* a moving target. Where Study 1’s amber rating reflected a capability gap on a verifiable sub-judgment, today’s open models close or exceed it; reliability there is largely a matter of time and access. But where the shortfall reflects a *compressed or contested human ceiling* (e.g. fluency, and, as Study 3 will show, subjective social annotation), even the most powerful models do not break through,




Fig. 4. Study 2. Summary-level Spearman of the most powerful open judges on SummEval, per aspect, against the reported G-Eval/GPT-4 bar (dashed). Bars above the dashed line indicate an open model that *surpasses* the reported proprietary-judge correlation.

because the limit lives in the human reference, not the model. This distinction, temporal capability gap vs. structural ceiling, is the lens we carry into HCC.

Takeaway. Study 1 says LLM judges are safe where the task is objective and verifiable, and Study 2 shows that frontier open models are now strong enough to match proprietary judges on exactly those verifiable tasks. But subjective, high-disagreement constructs (the kind that humans themselves contest) are precisely the tasks Study 1 flags as risky and Study 2 predicts scale will *not* fix. Study 3 then tests this prediction directly, on datasets with rich human label distributions.

5 STUDY 3: AN OPEN-PANEL META-EVALUATION ON NLI WITH RICH HUMAN LABELS

Study 1 shows LLM judges are most reliable where the human-agreement ceiling is high and verifiable; Study 2 shows frontier open models reach the established benchmark on the verifiable aspects of summarization. Study 3 tests the mechanism directly on *natural language inference* (NLI), the rare setting that ships *rich* human label distributions: ChaosNLI [25] re-annotates SNLI and MultiNLI items with ~100 human labels each, so the human-human agreement ceiling and the per-item disagreement are measured precisely rather than estimated from three annotators.

5.1 Method

We evaluate a panel of 6 open models served via an institutional vLLM API (OpenAI-compatible), spanning five families and 9B-754B parameters, run at temperature 0. Each model labels every item as a virtual annotator under the standard NLI instructions; we benchmark against the *human-human ceiling* (random-annotator-vs-majority accuracy with an annotator bootstrap 95% CI, plus pairwise agreement), and report accuracy, Cohen’s κ , and macro-F1 against the majority gold, a majority **jury** [27], and a calibrated contamination screen. The run comprised 6,000 judgments across 2

Table 3. Study 3 (NLI). Open panel vs the human-human ceiling. Ceiling = random-annotator-vs-majority accuracy; macro-F1 is the best model’s chance-corrected score.

Dataset	Construct	Human ceiling	Best model (acc)	macro-F1 (best)	Jury (acc)	Modal verdict
ChaosNLI-SNLI	NLI (3-way)	0.75	0.84 (gemma-4-31b)	0.79	0.86	Validated
ChaosNLI-MNLI	NLI (3-way)	0.64	0.70 (gemma-4-31b)	0.70	0.65	Promising

Table 4. Study 3 (NLI) disagreement-aware results. Ceiling with annotator-bootstrap 95% CI; best-model κ and macro-F1; panel accuracy on low- vs high-human-disagreement items; Jensen-Shannon panel-vs-human divergence.

Dataset	Ceiling [95% CI]	κ (best)	macro-F1 (best)	acc _{↓H} low-disagr.	acc _{↑H} high-disagr.	JS
ChaosNLI-SNLI	0.75 [0.74,0.77]	0.71	0.79	0.89	0.67	0.09
ChaosNLI-MNLI	0.64 [0.63,0.65]	0.53	0.70	0.68	0.55	0.17

NLI datasets with zero API errors and a 100% answer-parse rate. Dataset cards, prompts, and metric definitions are in Appendices A to E.

5.2 Results

Table 3 reports the panel against the human ceiling. Where humans agree only moderately (SNLI ceiling 0.75, MNLI 0.64), the open panel *matches or exceeds* them: the jury reaches 0.86 on SNLI ($>$ ceiling) and 5/6 models are **Validated**, with MNLI **Promising**. Chance-corrected agreement is substantial (best-model κ up to 0.71), so this is genuine competence, not base-rate learning.

5.3 Item-level disagreement governs reliability

The organizing finding is that errors concentrate on *contested* items. On both datasets, panel accuracy on low-human-disagreement (low-entropy) items far exceeds accuracy on high-disagreement items (0.79 vs 0.61; Table 4), and the panel’s label distribution diverges from the human distribution most on the most contested items. An item-level mixed-effects model of judge correctness (over all 6,000 judgments) makes this precise: the random-intercept SD is far larger across *items* (SD = 2.02) than across datasets (0.73) or models (between-model accuracy SD = 0.086), and with standardized predictors model scale contributes only modestly ($OR_{SD} = 1.21$). In short, *which item* is judged, its latent human (dis)agreement, predicts judge correctness better than which model judges it. A calibrated contamination screen (Appendix E) does not explain the pattern: it false-positives on high-recall models even for guaranteed-clean text, so we treat it as a conservative screen and rest the claim on the disagreement analysis. Juries raise accuracy but cannot lift agreement above the human ceiling on the most contested items.

6 DISCUSSION

6.1 What governs LLM-judge reliability

Across the three studies one principle organizes the evidence: an LLM judge is reliable to the degree that the task has a high, verifiable human-agreement ceiling, and it cannot exceed an agreement level that humans themselves do not reach. Study 3 localizes the mechanism at the *item* level. The random-intercept SD of judge correctness is far larger

across items ($SD = 2.02$) than across datasets (0.73) or models (between-model accuracy $SD = 0.086$); on *every* dataset, panel accuracy on low-human-disagreement items far exceeds that on high-disagreement items; and chance-corrected agreement (κ , macro-F1) collapses precisely where humans disagree. Model scale helps only modestly ($OR_{SD} = 1.21$), and a calibrated contamination screen argues the shortfall is not memorization. We are careful about the unit of analysis: with few datasets we do not identify a dataset-level coefficient (inconclusive, not null), so we locate our inferential claims at the well-powered item level, following standard practice for designs with few level-2 units.

6.2 Connection to measurement and disagreement scholarship

This result operationalizes a long-running argument that there is no single ground truth for subjective constructs [2, 25]: when judges are benchmarked against the human-human ceiling rather than a forced majority, their “failures” on contested items are revealed as properties of the construct’s contested human reference, not fixable model deficits. Reporting a single accuracy against a majority gold imposes consensus where the data shows disagreement; the honest object of evaluation is the human label *distribution*.

6.3 Implications for designing LLM-as-a-judge experiments

The evidence yields a concrete, reusable protocol. Any study using an LLM-as-a-judge should:

- (1) **Anchor to the human-human ceiling**, not a single gold; report the ceiling (computed from the annotator distribution at the dataset’s native annotator count) and the judge-vs-ceiling gap.
- (2) **Match the metric to the construct**: on imbalanced tasks report macro-F1 and per-class scores and prefer disagreement-aware evaluation against the full label distribution over forced-majority accuracy.
- (3) **Quantify uncertainty**: bootstrap confidence intervals for every comparison; do not assert “beats/ matches” from point estimates.
- (4) **Stratify by item difficulty**: because variance lives in items, report where the judge fails (high human-disagreement items), not only mean accuracy.
- (5) **Probe contamination, but calibrate the probe**: evaluate on post-cutoff data where possible, and treat surface-recall probes as conservative screens (they false-positive on high-recall models).
- (6) **Pin the setup**: fix model identifier, version, decoding (temperature, seed); report robustness to prompt paraphrase, since small models are prompt-sensitive on subjective tasks.
- (7) **Use juries for aggregate decisions** but not as a ceiling-breaker: panels improve stability yet do not lift a subjective task above its human ceiling.
- (8) **Pre-register and release**: register the construct, the human baseline, and the decision threshold, and release item-level labels and aggregates.

6.4 A decision framework

For annotation and evaluation tasks, three regimes follow from the human-ceiling principle. **Automate (with light audit)**: objective, verifiable judgments (entailment, factual consistency, instruction-following) where the human ceiling is high, exactly where Study 3’s NLI panel reaches or exceeds it. **Calibrate and audit**: judgments with a low or moderate ceiling, where a judge can post high accuracy by predicting base rates, so report chance-corrected agreement (κ) and audit the high-disagreement items. **Keep humans in the loop**: constructs on which human annotators themselves

disagree sharply, where no open judge can exceed a ceiling humans do not reach. A critical caution for the field: agreement on older public benchmarks can reflect memorization [4], and a single headline accuracy can mask base-rate learning, so calibrated, disagreement-aware, post-cutoff evaluation should be the norm. The application of these ideas to the subjective, interpretive tasks of Human-Centered Computing is the subject of a companion paper.

7 LIMITATIONS

The *human* validation here is inherited from prior work, not added by us: Study 3’s ground truth is ChaosNLI’s ~100 human labels per item, and the Study 1 review codes each study against the human baselines those studies report. The review verdict is our interpretation (only 29 of 215 report an extractable numeric agreement value); we report an independent re-coding as a reproducibility check (Krippendorff $\alpha = 0.87$, 100% within one level; Appendix D), which is LLM-assisted and therefore measures rubric reproducibility. All 215+55 references were web-verified and two unverifiable entries removed. By design we evaluate *open* models (the reproducible choice for most researchers) at temperature 0; Study 2 uses one summarization benchmark with integer scoring, and Study 3’s deep human-label data are available for NLI rather than every domain. Subjective social-computing annotation, where the human ceiling is contested, is the focus of a companion paper. Findings reflect models available through mid-2026.

8 ETHICS AND REPRODUCIBILITY

Ethics. The review analyzes published research, not human subjects, and Study 3 reuses already-public, de-identified NLI datasets under their licenses; this work was therefore not human-subjects research requiring IRB review, and we sought an exemption determination consistent with our institution’s policy. We release only item-level labels and aggregate statistics. One ethical point bears on the method itself: treating the *majority* human label as “ground truth” for the ceiling can encode majority bias and erase minority readings, which is why we report the full disagreement (soft labels, entropy) rather than only the majority and recommend disagreement-aware evaluation. The contamination screen is reported precisely so “validated” verdicts are not silently inflated by pretraining exposure. **Reproducibility.** Every number in this paper regenerates from released code: the OpenAlex, Semantic Scholar, and web search logs; the coded corpus; the figure and table generators; the SummEval harness and the Study-3 judging/scoring/contamination/mixed-effects harness, all keyed to documented model identifiers served through an institutional vLLM API at temperature 0 with idempotent, resumable checkpoints. The auto-generated statistics macros mean the manuscript cannot drift from the data. Repository: a public repository (URL omitted for anonymous review).

9 CONCLUSION

Open, locally-servable LLMs are now strong enough to serve as annotators on objective and verifiable tasks, where they meet or exceed established proprietary-judge correlations. But across 215 reviewed studies and a 6,000-judgment meta-evaluation on NLI datasets with rich human labels, reliability is bounded by the same thing for every model: the *human-human agreement ceiling*, and it degrades on exactly the items where humans themselves disagree. Scaling the judge helps only modestly, and contamination does not explain the pattern. The practical consequence is that an LLM-as-a-judge should be reported the way we report any measurement instrument, against a human reference, with chance-corrected and disagreement-aware metrics, with uncertainty, and with a stated operating range, rather than trusted as an oracle. The right question is never “can a machine judge?” but “can this open model agree with humans, on this construct, as well as humans agree with each other?”

REFERENCES

- [1] Argyle et al. 2023. Out of One, Many: Using Language Models to Simulate Human Samples. *arXiv preprint arXiv:2209.06899* (2023). arXiv:2209.06899
- [2] Lora Aroyo and Chris Welty. 2015. Truth Is a Lie: Crowd Truth and the Seven Myths of Human Annotation. In *AI Magazine*, Vol. 36. 15–24.
- [3] Bavaresco et al. 2024. LLMs instead of Human Judges? A Large Scale Empirical Study across 20 NLP Evaluation Tasks. *arXiv preprint arXiv:2406.18403* (2024). arXiv:2406.18403
- [4] Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FACt)*. 610–623.
- [5] Bisbee et al. 2023. Synthetic Replacements for Human Survey Data? The Perils of Large Language Models. *SocArXiv* (2023). <https://doi.org/10.31235/osf.io/5ecfa>
- [6] Calderon et al. 2025. The Alternative Annotator Test for LLM-as-a-Judge: How to Statistically Justify Replacing Human Annotators with LLMs. *arXiv preprint arXiv:2501.10970* (2025). arXiv:2501.10970
- [7] Chakrabarty et al. 2024. Art or Artifice? Large Language Models and the False Promise of Creativity. *arXiv preprint arXiv:2309.14556* (2024). arXiv:2309.14556
- [8] Chan et al. 2023. ChatEval: Towards Better LLM-based Evaluators through Multi-Agent Debate. *arXiv preprint arXiv:2308.07201* (2023). arXiv:2308.07201
- [9] Cheng et al. 2023. CoMPoS: Characterizing and Evaluating Caricature in LLM Simulations. In *EMNLP*. <https://doi.org/10.18653/v1/2023.emnlp-main.669>
- [10] Clarke and Dietz. 2024. LLM-based Relevance Assessment Still Can't Replace Human Relevance Assessment. *arXiv preprint arXiv:2412.17156* (2024). arXiv:2412.17156
- [11] Dai et al. 2023. LLM-in-the-loop: Leveraging Large Language Model for Thematic Analysis. In *EMNLP-F*. <https://doi.org/10.18653/v1/2023.findings-emnlp.669>
- [12] Davison et al. 2024. The Ethics of Using Generative AI for Qualitative Data Analysis. *Information Systems J.* (2024). <https://doi.org/10.1111/isj.12504>
- [13] Alexander R. Fabbri, Wojciech Kryściński, Bryan McCann, Caiming Xiong, Richard Socher, and Dragomir Radev. 2021. SummEval: Re-evaluating Summarization Evaluation. *Transactions of the Association for Computational Linguistics* 9 (2021), 391–409.
- [14] Gao et al. 2024. CollabCoder: A Lower-barrier, Rigorous Workflow for Inductive Collaborative Qualitative Analysis with Large Language Models. *arXiv preprint arXiv:2304.07366* (2024). arXiv:2304.07366
- [15] Gilardi et al. 2023. ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks. *arXiv preprint arXiv:2303.15056* (2023). arXiv:2303.15056
- [16] Gu et al. 2024. A Survey on LLM-as-a-Judge. *arXiv preprint arXiv:2411.15594* (2024). arXiv:2411.15594
- [17] Hashemi et al. 2024. LLM-Rubric: A Multidimensional, Calibrated Approach to Automated Evaluation of Natural Language Texts. *arXiv preprint arXiv:2501.00274* (2024). arXiv:2501.00274
- [18] James et al. 2026. Counting on Consensus: Selecting the Right Inter-annotator Agreement Metric for NLP Annotation and Evaluation. *arXiv preprint arXiv:2603.06865* (2026). arXiv:2603.06865
- [19] Kim et al. 2023. Prometheus: Inducing Fine-grained Evaluation Capability in Language Models. *arXiv preprint arXiv:2310.08491* (2023). arXiv:2310.08491
- [20] Kocmi and Federmann. 2023. GEMBA-MQM: Detecting Translation Quality Error Spans with GPT-4. *arXiv preprint arXiv:2310.13988* (2023). arXiv:2310.13988
- [21] Li et al. 2024. From Generation to Judgment: Opportunities and Challenges of LLM-as-a-judge. *arXiv preprint arXiv:2411.16594* (2024). arXiv:2411.16594
- [22] Li et al. 2024. LLMs-as-Judges: A Comprehensive Survey on LLM-based Evaluation Methods. *arXiv preprint arXiv:2412.05579* (2024). arXiv:2412.05579
- [23] Liu et al. 2023. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment. *arXiv preprint arXiv:2303.16634* (2023). arXiv:2303.16634
- [24] Matthew J. Page, Joanne E. McKenzie, Patrick M. Bossuyt, Isabelle Boutron, Tammy C. Hoffmann, Cynthia D. Mulrow, et al. 2021. The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ* 372 (2021), n71. <https://doi.org/10.1136/bmj.n71>
- [25] Barbara Plank. 2022. The "Problem" of Human Label Variation: On Ground Truth in Data, Modeling and Evaluation. *Proceedings of EMNLP* (2022), 10671–10682.
- [26] Schroeder et al. 2025. Large Language Models in Qualitative Research: Uses, Tensions, and Intentions. In *CHI*. <https://doi.org/10.1145/3706598.3713120>
- [27] Verga et al. 2024. Replacing Judges with Juries: Evaluating LLM Generations with a Panel of Diverse Models. *arXiv preprint arXiv:2404.18796* (2024). arXiv:2404.18796
- [28] Wang et al. 2023. PandaLM: An Automatic Evaluation Benchmark for LLM Instruction Tuning Optimization. *arXiv preprint arXiv:2306.05087* (2023). arXiv:2306.05087
- [29] Wang et al. 2024. Large Language Models that Replace Human Participants can Harmfully Misportray and Flatten Identity Groups. *arXiv preprint arXiv:2402.01908* (2024). arXiv:2402.01908
- [30] Wei et al. 2024. Long-form Factuality in Large Language Models. *arXiv preprint arXiv:2403.18802* (2024). arXiv:2403.18802
- [31] Xiao et al. 2023. Supporting Qualitative Analysis with Large Language Models: Combining Codebook with GPT-3 for Deductive Coding. *arXiv preprint arXiv:2304.10548* (2023). arXiv:2304.10548
- [32] Zheng et al. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *arXiv preprint arXiv:2306.05685* (2023). arXiv:2306.05685
- [33] Zhu et al. 2023. JudgeLM: Fine-tuned Large Language Models are Scalable Judges. *arXiv preprint arXiv:2310.17631* (2023). arXiv:2310.17631

A REVIEW PROTOCOL AND RELIABILITY RUBRIC

Identification. Both reviews follow PRISMA 2020. The primary bibliographic source is the **OpenAlex** API; we cross-check with the **Semantic Scholar** Graph API and a structured web “**deep-research**” arm. Study 1a used a 59-query web protocol plus venue-scoped sweeps; Study 2 used a 25-query HCC/HCI/qualitative concept-block protocol (the same query strings were issued against every source so the arms are comparable). Concept blocks combine an LLM term (“LLM”, “large language model”, “GPT-4”, “ChatGPT”, “Gemini”), an evaluation/agreement term (“LLM-as-a-judge”, “automatic evaluation”, “annotation”, “inter-annotator agreement”, “human evaluation”), and, for Study 2, an HCC method term (“qualitative coding”, “thematic analysis”, “content analysis”, “crowdsourcing”, “participant simulation”, “persona”, “usability evaluation”, “design critique”). Records were deduplicated by DOI and normalized title.

Screening and eligibility. *Inclusion:* peer-reviewed or preprint studies that empirically compare LLM judgments to human judgments on the same items and report (or allow computing) an agreement statistic. *Exclusion:* non-empirical position papers, studies with no human baseline, studies that use LLMs only to generate (not judge) content, and off-topic retrievals. Title/abstract screening was followed by full-text eligibility assessment; per-stage counts appear in the PRISMA flow figures (Figure 1).

Reliability verdict (five-level rubric). Each study is coded against the *human-human agreement ceiling* for its task:

- **1 Validated:** LLM-human agreement is at or above the human-human ceiling (substitution is empirically justified).
- **2 Promising:** agreement approaches the ceiling under specific, reported conditions (model, prompt, aggregation); conditional substitution.
- **3 Mixed:** results are task- or setup-dependent and inconsistent across studies.
- **4 Caution:** agreement is consistently and materially below the ceiling.
- **5 Unreliable:** agreement is near chance or dominated by bias (position, verbosity, self-preference).

We report a reliability score = 6 – verdict and synthesize within metric families (e.g. Cohen’s κ , Krippendorff’s α , Spearman ρ , accuracy) rather than pooling incommensurable effect sizes. The full coded corpora (Appendix F) are released for re-coding.

B STUDY 1B: G-EVAL REPLICATION DETAILS

We replicate the summary-level G-Eval protocol [23] on SummEval [13] (100 source articles $\times$ 16 system summaries = 1600 summaries, each with expert means for four aspects). Each (summary, aspect) pair is scored 1 to 5 by the judge with the aspect criterion below; we then compute the **summary-level Spearman ρ** (per source document, correlate the judge’s 16 scores with the human means, averaged over documents, dropping documents with no score variance) and compare it to the reported G-Eval-4/GPT-4 correlations. Judges are the four most powerful open models in the available open-model catalogue (glm-5.1-fp8 754B, deepseek-v4-pro, gptoss-120b, qwen3.6-35b-a3b-fp8), decoded at temperature 0. Aspect criteria (verbatim):

Coherence (1-5): the collective quality of all sentences; the summary should be well-structured and well-organized, building from sentence to sentence into a coherent body of information, not just a heap of related information.

Consistency (1-5): factual alignment between the summary and the source; a consistent summary contains only statements entailed by the source document; hallucinated or unsupported facts are penalized.

Fluency (1-5): the quality of individual sentences (grammar, capitalization, readability).

Relevance (1-5): selection of the most important content from the source; redundancy and unimportant content are penalized.

The judge receives the source article and the summary and returns only a JSON object {"score": <1-5>} (a larger token budget and a retry handle reasoning-model outputs). In total Study 1b contributes 25,600 judgments.

C STUDY 3: DATASETS, JUDGING PROMPTS, AND METRICS

Datasets. Table 5 summarizes the six datasets. Each is sampled to $n=500$ items with a clear majority label (ties dropped); we use the per-item annotator distributions to compute the human ceiling.

Table 5. Study 3 dataset cards. Ceiling = random-annotator-vs-majority accuracy. Contam. = max verbatim-recall tier across probed judges.

Dataset	Task	Source	Annot./item	Ceiling	Contam.
ChaosNLI-SNLI	NLI (3-way)	SNLI + ChaosNLI relabels	~100	0.75	MEDIUM
ChaosNLI-MNLI	NLI (3-way)	MultiNLI + ChaosNLI relabels	~100	0.64	MEDIUM
GoEmotions	sentiment (4-way, authors' grouping)	Reddit (Google)	~3-5	0.78	MEDIUM
SBIC	offensiveness (binary)	Social Bias Frames	~3	0.91	MEDIUM
HateXplain	hate speech (3-way)	Twitter/Gab	3	0.83	MEDIUM
MultiPICO	irony (binary)	Reddit/Twitter (2024)	3-8	0.82	LOW/MED

Judging prompts. Each model judges every item as a virtual annotator under instructions paralleling the human guidelines, returning a JSON label on the final line (parsed with a surface-form-to-code map). The exact user prompts:

NLI: Premise: {premise} / Hypothesis: {hypothesis} / Does the premise ENTAIL the hypothesis, CONTRADICT it, or is it NEUTRAL? -> {"label": "entailment"|"neutral"|"contradiction"}

Hate: Post: {text} / Classify: HATE SPEECH (attacks/dehumanizes a group based on identity), OFFENSIVE (abusive but not group-targeted), or NORMAL. -> {"label": "hate speech"|"offensive"|"normal"}

Emotion: Comment: {text} / Overall sentiment: POSITIVE, NEGATIVE, AMBIGUOUS (mixed/uncertain), or NEUTRAL. -> {"label": "positive"|"negative"|"ambiguous"|"neutral"}

Offensive: Post: {text} / Is this post offensive (including potentially offensive)? -> {"label": "offensive"|"not offensive"}

Irony: {post+reply} / Is the REPLY ironic/sarcastic (conveys a meaning different from or opposite to its literal content), or not ironic? -> {"label": "ironic"|"not ironic"}

Metrics. Let an item have A human annotations with majority label y^* . **Human ceiling** is the random-annotator-vs-majority accuracy (the expected agreement of one held-out human with the majority), averaged over items; we also report the pairwise human-human agreement. **Accuracy**, **Cohen's κ** , and **macro-F1** are computed against y^* . The **alt-test** winning rate [6] is the fraction of annotators for whom replacing that human with the LLM does not reduce agreement with the remaining humans (a count-based leave-one-annotator-out test); a winning rate ≥ 0.5 is the bar for claiming the LLM can “replace” a human. The **jury** is the per-item majority vote of the model panel [27]. **95% confidence intervals** are nonparametric bootstrap (1000 item resamples). **Power:** at $n=500$ the accuracy CI half-width is $\leq \pm 4.4\%$ and the design has $> 80\%$ power to detect a 0.10 gap from the ceiling. **Contamination probe:** the judge is given the dataset name and a prefix ($\approx$ one third) of an item's text and asked to reproduce the remainder verbatim; we report the mean normalized sequence overlap and flag HIGH (> 0.6 mean or $> 20\%$ near-verbatim), MEDIUM (> 0.35

mean), or LOW. Only item-level labels and aggregates are released (no raw text); the runs recorded *zero* API errors and a 100% answer-parse rate across all 18000 judgments.

Inferential model. The mixed-effects analysis models each judgment’s correctness ($1[\text{judge} = y^*]$) with fixed effects for task subjectivity (1–ceiling, dataset level) and model scale ($\log_{10}$ parameters), reported as a Bayesian logistic GLMM with random intercepts for dataset and item and, for clean inference, as a logistic GLM with item-cluster-robust standard errors; the Benjamini-Hochberg procedure controls the false-discovery rate across the 36 judge $\times$ dataset ceiling-gap tests. A naive two-predictor version of this model overstates a “subjectivity” effect; once construct granularity and label base-rate skew are added and item-cluster-robust standard errors are used, subjectivity is not significant (OR = 0.74, $p = 0.45$) and the robust predictors are base-rate skew, scale, and granularity. A dataset-level meta-regression ($n=6$) is reported with leave-one-dataset-out bounds.

D INTER-CODER RELIABILITY

Because the reliability verdict and the role/stance codes embed judgement, we re-coded a held-out random sample of each corpus with two *independent* coders who saw only the per-study summaries and the rubric (no original codes, no other context). For Study 1a ($n=54$, 25%) the verdict reproduced at Krippendorff $\alpha = 0.87$ (ordinal, three coders) with 100% of cases within one level. For Study 2 ($n=16$, 30%) the verdict gave $\alpha = 0.85$, the augment/replace/simulate role $\alpha = 0.80$, and the epistemic stance $\alpha = 0.59$ (the least reliable code). *Caveat:* the re-coders are LLM-assisted, like the primary pipeline, so these figures measure rubric *reproducibility* under independent application, not human inter-rater agreement; the moderate stance α is why we treat stance-based claims as indicative.

E ROBUSTNESS AND VALIDITY CHECKS

Prompt robustness. We re-ran a subset of Study 3 (three models $\times$ six datasets, $n=120$) under three semantically equivalent prompt paraphrases per task. *Accuracy* is fairly prompt-stable (mean SD across paraphrases = 0.03), but the discrete five-level *verdict* is exactly stable in only 10/18 cells: most instability is a one-level flip for cells sitting near a band boundary, and the smaller model is the most sensitive (worst-case acc SD 0.14), the larger models the least. We therefore report verdicts as robust to *within one level* rather than exact, and flag small-model verdicts on subjective tasks as prompt-sensitive. **Contamination clean control.** We ran the verbatim-recall probe on freshly authored, post-cutoff sentences that no model can have seen. The result is a cautionary one: the high-recall model still scores 0.45 mean overlap (MEDIUM) on this guaranteed-clean text, i.e. it confabulates plausible continuations regardless of memorization, while the low-recall model scores LOW. The probe thus has a high false-positive rate for high-recall models and *cannot* be read as a reliable leakage test; the uniform MEDIUM tiers in Table 3 are largely an artifact of generative fluency. We retain the probe only as a conservative, low-recall-model signal and base our substantive claims on the disagreement analysis instead. **Annotator-subgroup check.** On MultiPICo (which carries annotator demographics), the panel’s agreement with the female- vs male-annotator majority differs by only 0.01, i.e. we did not detect a large majority-group bias in this corpus, though demographic metadata is sparse and this check is limited. **Ceiling vs annotator count.** The random-annotator-vs-majority ceiling depends on how many annotators an item has: subsampling ChaosNLI-SNLI (~ 100 annotators) to $A=3$ raises its estimated ceiling from 0.75 to 0.81 (fewer annotators yield more unanimous majorities). We therefore compute each dataset’s ceiling at its *native* annotator count and caution that ceilings are not directly comparable across datasets with very different A ; the disagreement-aware analyses (entropy, κ) do not depend on this estimator.

F STUDY 1 CORPUS

Table 6. Full extracted-studies corpus (215 included studies). Agreement is the headline coefficient/proportion where reported.

Key	Study	Domain	Task	Metric	Agree	Verdict
mirzakhmedova2024argument	Mirzakhmedova et al. (2024)	Argument Quality	pointwise	pct_agreement	-	Mixed
tong2024codejudge	Tong et al. (2024)	Code & SW Eng	classify	accuracy	-	Promising
watson2024aicodereview	Watson et al. (2024)	Code & SW Eng	classify	pct_agreement	-	Mixed
codespec2025	Jin et al. (2025)	Code & SW Eng	verify	accuracy	-	Caution
wang2025setesting	Wang et al. (2025)	Code & SW Eng	pairwise	pct_agreement	-	Mixed
zhuo2025codejudgeeff	Wang et al. (2025)	Code & SW Eng	pointwise	kappa	-	Mixed
biasloop2026	Zhao et al. (2026)	Code & SW Eng	pairwise	pct_agreement	-	Caution
overcorrect2026	Jin et al. (2026)	Code & SW Eng	classify	accuracy	-	Caution
hatespeechprobe2023	Roy et al. (2023)	Content Moderation & Toxicity	classify	f1	-	Mixed
li2023hotchatgpt	Li et al. (2023)	Content Moderation & Toxicity	classify	pct_agreement	-	Mixed
nguyen2024crisismisinfo	Nguyen & Rudra (2024)	Content Moderation & Toxicity	classify	f1	-	Mixed
hatespeechanno2025	authors (2025)	Content Moderation & Toxicity	classify	kappa	-	Caution
icwsm2025hatebias	Giorgi et al. (2025)	Content Moderation & Toxicity	classify	kappa	-	Caution
xie2023nextchapter	Xie et al. (2023)	Creative Writing	pointwise	pct_agreement	-	Mixed
chakrabarty2024artartifice	Chakrabarty et al. (2024)	Creative Writing	pointwise	pearson	0.10	Unreliable
ismayilzada2024creative	Ismayilzada et al. (2024)	Creative Writing	pointwise	pct_agreement	-	Caution
creativeevalmdpi2025	Kim et al. (2025)	Creative Writing	pointwise	spearman	-	Mixed
haase2025creativitypeaked	Haase et al. (2025)	Creative Writing	pointwise	pct_agreement	-	Caution
litbench2025	Fein et al. (2025)	Creative Writing	pairwise	pct_agreement	0.78	Mixed
lin2023l1meval	Lin & Chen (2023)	Dialogue	pointwise	spearman	-	Promising
hackl2023reliablerater	Hackl et al. (2023)	Education & Grading	pointwise	icc	-	Promising
khademi2023bardassessment	Khademi (2023)	Education & Grading	pointwise	pct_agreement	-	Mixed
kortemeyer2023physics	Kortemeyer (2023)	Education & Grading	pointwise	pct_agreement	-	Mixed
mizumoto2023aes	Mizumoto & Eguchi (2023)	Education & Grading	pointwise	qwk	0.74	Promising
yancey2023cefr	Yancey et al. (2023)	Education & Grading	pointwise	qwk	-	Mixed
aescomparative2024	Lee et al. (2024)	Education & Grading	pointwise	qwk	0.57	Mixed
floden2024gradingexams	Flodén (2024)	Education & Grading	pointwise	pct_agreement	-	Mixed
grevisse2024asagmed	Grévisse (2024)	Education & Grading	classify	pct_agreement	-	Promising
hou2024classroom	Hou et al. (2024)	Education & Grading	pointwise	icc	-	Mixed
k12shortanswer2024	Henkel et al. (2024)	Education & Grading	classify	kappa	0.70	Promising
koraishi2024langassess	Koraishi (2024)	Education & Grading	pointwise	pct_agreement	-	Mixed
quah2024dentalaes	Quah et al. (2024)	Education & Grading	pointwise	icc	-	Mixed
wang2024beyondagreement	Wang et al. (2024)	Education & Grading	pointwise	qwk	-	Mixed
asagfair2025	Henkel et al. (2025)	Education & Grading	pointwise	kappa	-	Mixed
sefl2025feedback	authors (2025)	Education & Grading	pointwise	f1	-	Promising
studentsjudge2025	Banihashem et al. (2025)	Education & Grading	pointwise	pct_agreement	-	Mixed
k12sci2026	He et al. (2026)	Education & Grading	pointwise	pct_agreement	-	Promising
claimmatch2023	Choi et al. (2023)	Fact-Checking	classify	pct_agreement	-	Promising
ni2024afacta	Ni et al. (2024)	Fact-Checking	classify	pct_agreement	-	Promising
politicaltruths2024	Chatrath et al. (2024)	Fact-Checking	classify	pct_agreement	-	Mixed
factsopinions2025	Saju et al. (2025)	Fact-Checking	classify	accuracy	-	Caution
kaikaus2023financecorpora	Kaikaus et al. (2023)	Finance	classify	accuracy	-	Mixed
bizfinbench2025	Lu et al. (2025)	Finance	pointwise	pct_agreement	-	Mixed
trident2025	Hui et al. (2025)	Finance	classify	pct_agreement	-	Mixed
chan2023chateval	Chan et al. (2023)	General NLG	pairwise	pct_agreement	-	Mixed
chiang2023alt	Chiang & Lee (2023)	General NLG	pointwise	pct_agreement	-	Mixed
dubois2023alpacafarm	Dubois et al. (2023)	General NLG	pairwise	winrate_corr	0.98	Promising
hsu2023figurecaption	Hsu et al. (2023)	General NLG	pointwise	spearman	-	Promising
kim2023prometheus	Kim et al. (2023)	General NLG	pointwise	pearson	0.90	Promising
lai2023styletransfer	Lai et al. (2023)	General NLG	pointwise	spearman	-	Mixed

	Key	Study	Domain	Task	Metric	Agree	Verdict
781							
782	li2023autoj	Li et al. (2023)	General NLG	pairwise	pct_agreement	-	Mixed
783	wang2023chatgptnl	Wang et al. (2023)	General NLG	pointwise	spearman	-	Mixed
784	wang2023pandalm	Wang et al. (2023)	General NLG	pairwise	accuracy	-	Promising
785	xu2023instructscore	Xu et al. (2023)	General NLG	pointwise	pearson	-	Promising
786	zhang2023widerdeeper	Zhang et al. (2023)	General NLG	pairwise	pct_agreement	-	Mixed
787	zheng2023mtbench	Zheng et al. (2023)	General NLG	pairwise	pct_agreement	0.85	Validated
788	chiang2024arena	Chiang et al. (2024)	General NLG	pairwise	pct_agreement	-	Validated
789	doostmohammadi2024autoeval	Doostmohammadi et al. (2024)	General NLG	pointwise	pct_agreement	-	Mixed
790	dubois2024lc	Dubois et al. (2024)	General NLG	pairwise	winrate_corr	0.98	Promising
791	hashemi2024llmrubric	Hashemi et al. (2024)	General NLG	pointwise	pearson	-	Promising
792	huang2024nlexpl	Huang et al. (2024)	General NLG	pointwise	pct_agreement	-	Mixed
793	kim2024prometheus2	Kim et al. (2024)	General NLG	pointwise	pearson	-	Promising
794	li2024arenahard	Li et al. (2024)	General NLG	pairwise	pct_agreement	0.89	Promising
795	lin2024wildbench	Lin et al. (2024)	General NLG	pairwise	winrate_corr	0.98	Promising
796	verga2024poll	Verga et al. (2024)	General NLG	pairwise	pct_agreement	-	Promising
797	faggioli2023perspectives	Faggioli et al. (2023)	IR & Relevance	classify	kappa	-	Promising
798	alaofi2024fooledrelevant	Alaofi et al. (2024)	IR & Relevance	classify	pct_agreement	-	Caution
799	clarke2024cantreple	Clarke & Dietz (2024)	IR & Relevance	classify	kappa	-	Caution
800	thomas2024relevance	Thomas et al. (2024)	IR & Relevance	classify	kappa	-	Validated
801	upadhyay2024largescale	Upadhyay et al. (2024)	IR & Relevance	classify	kappa	-	Promising
802	zhuang2024beyondyesno	Zhuang et al. (2024)	IR & Relevance	ranking	kendall	-	Promising
803	arabzadeh2025benchrelevance	Arabzadeh & Clarke (2025)	IR & Relevance	classify	kappa	-	Mixed
804	arabzadeh2025promptsens	Arabzadeh & Clarke (2025)	IR & Relevance	classify	kappa	-	Mixed
805	frobe2025assessorsagree	Frobe et al. (2025)	IR & Relevance	classify	kappa	-	Caution
806	li2023coannotating	Li et al. (2023)	Information Extraction (NER/NLI)	classify	pct_agreement	-	Promising
807	humanllmcollab2024	Wang et al. (2024)	Information Extraction (NER/NLI)	classify	pct_agreement	-	Promising
808	kholodna2024activelearning	Kholodna et al. (2024)	Information Extraction (NER/NLI)	classify	f1	-	Promising
809	kim2024meganno	Kim et al. (2024)	Information Extraction (NER/NLI)	classify	f1	-	Promising
810	nasution2024chatgptlabel	Nasution & Onan (2024)	Information Extraction (NER/NLI)	classify	accuracy	-	Mixed
811	ner2024augment	Naraki et al. (2024)	Information Extraction (NER/NLI)	classify	f1	-	Promising
812	nli2024labeldist	Chen et al. (2024)	Information Extraction (NER/NLI)	classify	pearson	-	Promising
813	ronningstad2024gptannotator	Ronningstad et al. (2024)	Information Extraction (NER/NLI)	classify	f1	-	Mixed
814	qiu2025labelensemble	Qiu et al. (2025)	Information Extraction (NER/NLI)	classify	accuracy	-	Promising
815	zhou2023ifeval	Zhou et al. (2023)	Instruction Following	verify	accuracy	-	Validated
816	zeng2024llmbar	Zeng et al. (2024)	Instruction Following	pairwise	accuracy	-	Mixed
817	ifcritic2025	Wen et al. (2025)	Instruction Following	classify	pct_agreement	-	Promising
818	legaldocrec2025	Pradhan et al. (2025)	Legal	ranking	spearman	0.73	Promising
819	lemaj2025	Enguehard et al. (2025)	Legal	pointwise	pct_agreement	-	Mixed
820	kocmi2023gemba	Kocmi & Federmann (2023)	Machine Translation	pointwise	kendall	-	Promising
821	kocmi2023gembamqm	Kocmi & Federmann (2023)	Machine Translation	verify	pct_agreement	0.96	Validated
822	lu2023erroranalysis	Lu et al. (2023)	Machine Translation	verify	kendall	-	Promising
823	sun2023radimpressions	Sun et al. (2023)	Medicine & Clinical	pointwise	pct_agreement	-	Caution
824	tang2023medevidence	Tang et al. (2023)	Medicine & Clinical	pointwise	pct_agreement	-	Caution
825	brain2024radgpt	Mitsuyama et al. (2024)	Medicine & Clinical	classify	accuracy	0.94	Mixed
826	brake2024clinicalnote	Brake & Schaaf (2024)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
827	wang2024clinicalevent	Wang et al. (2024)	Medicine & Clinical	classify	pct_agreement	-	Promising
828	chung2025verifact	Chung et al. (2025)	Medicine & Clinical	verify	pct_agreement	-	Mixed
829	clinstruct2025	Brake et al. (2025)	Medicine & Clinical	pointwise	icc	0.82	Mixed
830	croxford2025autoeval	Croxford et al. (2025)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
831	croxford2025clinsummjudge	Croxford et al. (2025)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
832	healthbench2025	Arora et al. (OpenAI) (2025)	Medicine & Clinical	pointwise	pct_agreement	-	Promising
833	kim2025emulate	Kim & Kim (2025)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
834	llmevalmed2025	Zhang et al. (2025)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
835	szymanski2025limits	Szymanski et al. (2025)	Medicine & Clinical	pairwise	pct_agreement	0.66	Caution

Key	Study	Domain	Task	Metric	Agree	Verdict
clindialog2026	Sun et al. (2026)	Medicine & Clinical	pointwise	pct_agreement	-	Mixed
frenchmedqa2026	Belmadani et al. (2026)	Medicine & Clinical	pointwise	kappa	-	Mixed
medjudge2026scoping	Li et al. (2026)	Medicine & Clinical	review	nan	-	Caution
counselbench2025	Li et al. (2025)	Mental Health	pointwise	pct_agreement	-	Caution
hua2025mhscoping	Hua et al. (2025)	Mental Health	review	nan	-	Caution
mhtrust2025	Badawi et al. (2025)	Mental Health	pointwise	icc	-	Mixed
mhsupport2026	Badawi et al. (2026)	Mental Health	pointwise	icc	-	Mixed
liu2023trustworthy	Liu et al. (2023)	Methodology & Bias	survey	nan	-	Mixed
saito2023verbositybias	Saito et al. (2023)	Methodology & Bias	pairwise	pct_agreement	-	Caution
wang2023shepherd	Wang et al. (2023)	Methodology & Bias	pointwise	pct_agreement	-	Promising
zhu2023judge1m	Zhu et al. (2023)	Methodology & Bias	pairwise	pct_agreement	-	Promising
bavaresco2024judgebench	Bavaresco et al. (2024)	Methodology & Bias	classify	kappa	0.28	Mixed
chen2024humansor1lms	Chen et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Caution
desmond2024eval1lm	Desmond et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
fan2024sedareval	Fan et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Promising
feuer2024style	Feuer et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Caution
gu2024survey	Gu et al. (2024)	Methodology & Bias	survey	nan	-	Mixed
jung2024trustescalate	Jung et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Promising
justice2024prejudice	Ye et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Caution
ke2024critiquellm	Ke et al. (2024)	Methodology & Bias	pointwise	spearman	-	Mixed
koo2024cobbler	Koo et al. (2024)	Methodology & Bias	pairwise	winrate_corr	0.44	Caution
li2024compsurvey	Li et al. (2024)	Methodology & Bias	survey	nan	-	Mixed
li2024gen2judge	Li et al. (2024)	Methodology & Bias	survey	nan	-	Mixed
li2024splitmerge	Li et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Mixed
liu2024pairwisepref	Liu et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Promising
pan2024hcdjudge	Pan et al. (2024)	Methodology & Bias	pointwise	nan	-	Mixed
panickssery2024selfpref	Panickssery et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Caution
raina2024robust	Raina et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Unreliable
shankar2024validators	Shankar et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
stureborg2024inconsistent	Stureborg et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Caution
thakur2024judgingjudges	Thakur et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Mixed
wataoka2024selfpref	Wataoka et al. (2024)	Methodology & Bias	pairwise	pct_agreement	-	Caution
zhu2024nationalitybias	Zhu et al. (2024)	Methodology & Bias	pointwise	pct_agreement	-	Caution
calderon2025a1ttest	Calderon et al. (2025)	Methodology & Bias	classify	pct_agreement	-	Promising
gao2025nlgsurvey	Gao et al. (2025)	Methodology & Bias	survey	nan	-	Mixed
gao2025reranking	Gao et al. (2025)	Methodology & Bias	ranking	kendall	-	Mixed
hu2025trainjudge	Hu et al. (2025)	Methodology & Bias	pointwise	pct_agreement	-	Promising
huang2025empiricaljudge	Huang et al. (2025)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
judgesverdict2025	Han et al. (2025)	Methodology & Bias	pointwise	kappa	-	Promising
li2025prefleakage	Li et al. (2025)	Methodology & Bias	pairwise	pct_agreement	-	Caution
murugadoss2025evaltheeval	Murugadoss et al. (2025)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
voutsas2025biaseddesign	Voutsas et al. (2025)	Methodology & Bias	classify	pct_agreement	-	Caution
yamauchi2025designchoices	Yamauchi et al. (2025)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
imrit2026iaa	James et al. (2026)	Methodology & Bias	review	krippendorff_alpha	-	Mixed
temperature2026	Li et al. (2026)	Methodology & Bias	pointwise	pct_agreement	-	Mixed
khondaker2023gptaraeval	Khondaker et al. (2023)	Multilingual	classify	f1	-	Caution
mmeval2024	Son et al. (2024)	Multilingual	pairwise	accuracy	-	Mixed
multiljudge2025	Fu et al. (2025)	Multilingual	pointwise	pct_agreement	-	Caution
reliablemulti2026	Zubiaga et al. (2026)	Multilingual	pointwise	pct_agreement	-	Caution
chen2024m1lmjudge	Chen et al. (2024)	Multimodal (Vision-Language)	pairwise	pct_agreement	0.79	Mixed
duan2024uimockup	Duan et al. (2024)	Multimodal (Vision-Language)	pointwise	pct_agreement	-	Mixed
lee2024promvision	Lee et al. (2024)	Multimodal (Vision-Language)	pointwise	pearson	-	Mixed
lu2024wildvision	Lu et al. (2024)	Multimodal (Vision-Language)	pairwise	pct_agreement	-	Promising
yasunaga2025mmrewardbench	Yasunaga et al. (2025)	Multimodal (Vision-Language)	pairwise	accuracy	0.72	Caution

	Key	Study	Domain	Task	Metric	Agree	Verdict
885							
886							
887							
888							
889							
890							
891							
892							
893							
894							
895							
896							
897							
898							
899							
900							
901							
902							
903							
904							
905							
906							
907							
908							
909							
910							
911							
912							
913							
914							
915							
916							
917							
918							
919							
920							
921							
922							
923							
924							
925							
926							
927							
928							
929							
930							
931							
932							
933							
934							
935							
936							

Key	Study	Domain	Task	Metric	Agree	Verdict
zhang2024newssumm	Zhang et al. (2024)	Summarization	pointwise	pct_agreement	-	Promising
wu2023hpsv2	Wu et al. (2023)	Text-to-Image	pointwise	accuracy	-	Promising
lin2024vqascore	Lin et al. (2024)	Text-to-Image	pointwise	pearson	-	Promising